

# Docker From The Ground Up Part 4

## Kubernetes 101

A whistle-stop tour of Kubernetes architecture and control/data planes.

Docker Birmingham Meetup 12/12/2018



kubernetes

**Thanks to our Sponsor!**

**BlackCat /**



**Thanks to Katacoda**

**For providing a great learning environment/demos**

Powered by  **Kata<oda**



# \$ whoami



**Shaun McLernon**

shaun@codesome.tech

*Independent software engineer with 20+ years in development roles, and 5 years in "devops" space.*

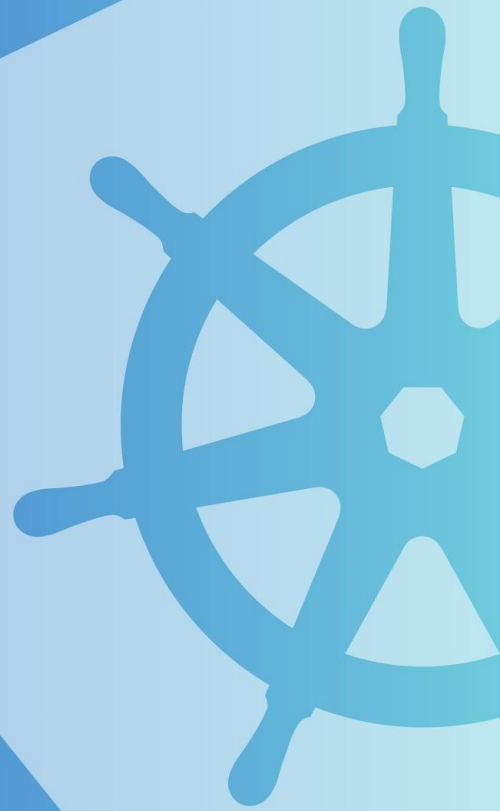
Github: @smclernon

Docker Birmingham Slack: @shaun



kubernetes

# Recap





# Requirements for Dev -> Production

- Multi-machine
- Discovery and Naming
- Scaling
- Multiple users
- Failure tolerance and recovery
- Monitoring
- Logging
- High availability
- Deployment lifecycle
- Load balancing
- etc, etc




# Alternatives



Apache  
**MESOS**<sup>™</sup>



HashiCorp  
**Nomad**



**AWS ECS**








Image ref: <https://www.cncf.io/blog/2018/08/29/cncf-survey-use-of-cloud-native-technologies-in-production-has-grown-over-200-percent/>

# Poll: Developers/Operators spread?



# Kubernetes

## K8s

Greek for "Helmsman" < person who steers a ship

K8s is an abbreviation derived by replacing the 8 letters “ubernete” with “8”.



Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.



# What does that really mean?

Kubernetes is a "Container Orchestrator"

- Scheduling: put containers on machines
  - by resource needs (CPU, memory)
  - by affinity requirements (put X near Y)
  - by labels put X on a "GPU" machines
- Replication: run N copies
- Discovery: find peers and services in other containers
- Recovery: automatically from failure
  - Health checking
  - Application failures
  - Machine failures
- Inspection: tell me what is happening
  - Basic monitoring
  - Logging



# Where did it come from?

- Created by Joe Beda, Craig McLuckie (Heptio) and Brendan Burns (Microsoft)
- Based on ideas proven at Google for a decade and a half of experience that Google has with running production workloads at scale. \*
- Project was open-sourced in 2014.
- Google has since donated to the Cloud Native Computing Foundation (CNCF).

\* [Large-scale cluster management at Google with Borg](#)



# Design Principles

- Declarative > imperative
- Control loops
- Simple > complex
- Modularity
- Legacy compatible
- Network-centric
- Labels > hierarchy
- Cattle > pets
- Open > closed



# Kubernetes Architecture





# Architecture overview

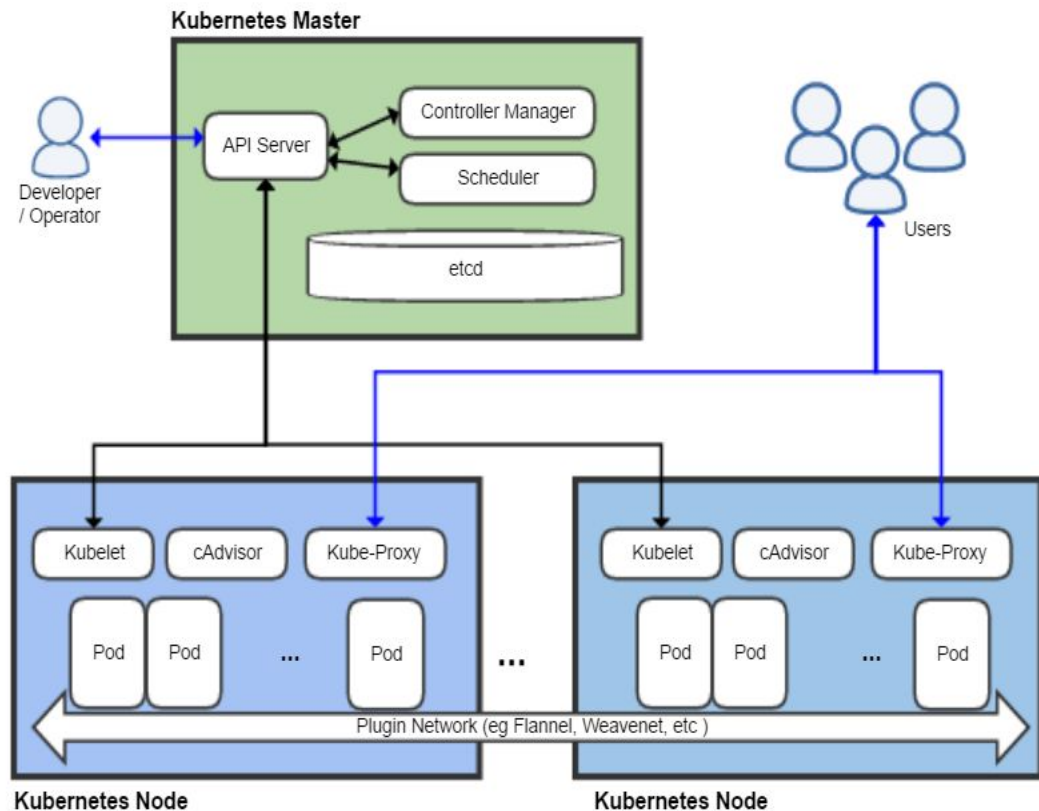
The machines that make up a Kubernetes cluster are called **nodes**.

Nodes in a Kubernetes cluster may be physical (bare metal), or virtual (vm's).

Two types of nodes:

**Masters**, which make up the **Control Plane**, acts as the “brains” of the cluster.

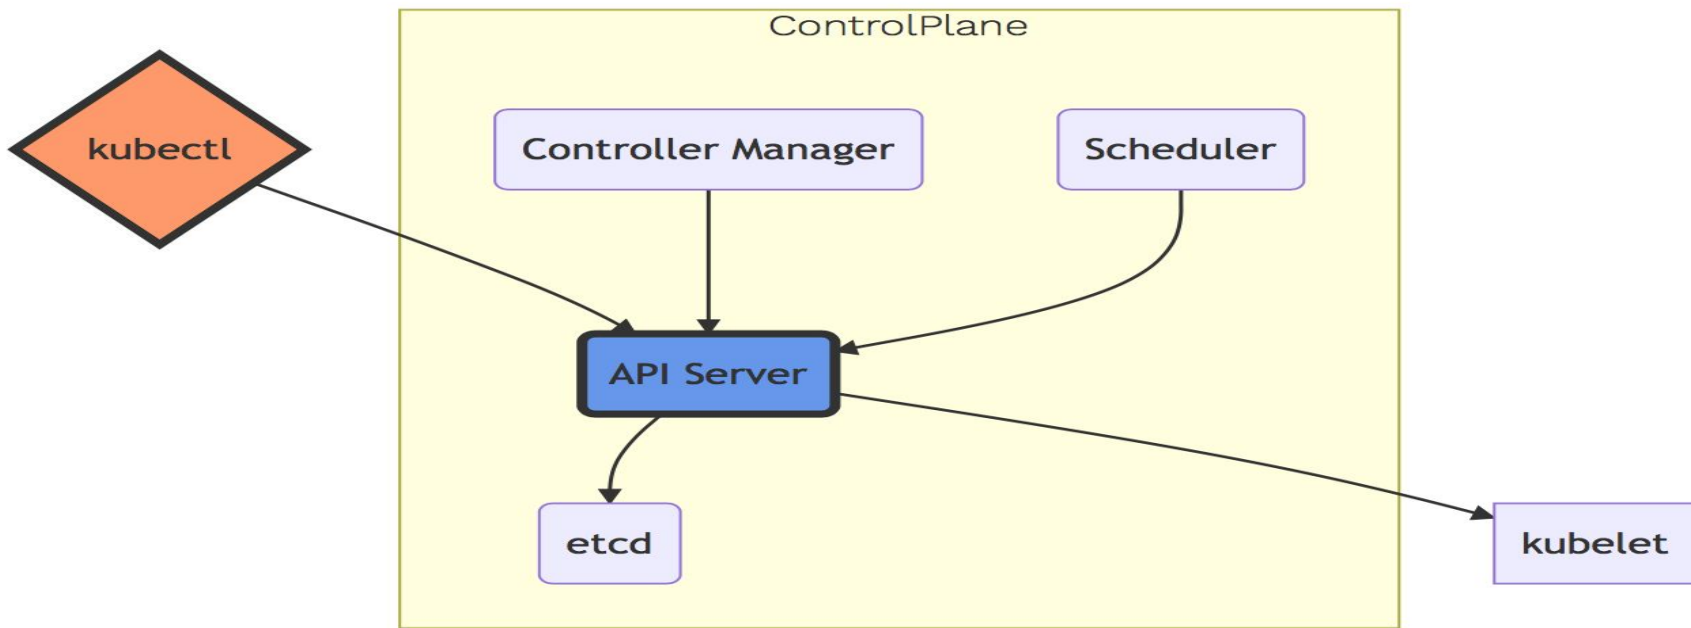
**Workers**, which make up the **Data Plane**, runs the actual container images (via pods).





# Master components (Control plane)

# Master components



- One or more API Servers: Entry point for REST / kubectl
- etcd: Distributed key/value store
- Controller-manager: Always evaluating current vs desired state
- Scheduler: Schedules pods to worker nodes

# API server



The kube API server exposes the Kubernetes REST API. It can easily scale horizontally as it is stateless and stores all data in the etcd cluster.

All clients, including nodes, users and other applications interact with Kubernetes strictly through the API Server.

# etcd



Etcd is a highly reliable distributed datastore. Kubernetes uses it to store the entire cluster state.

In small, transient cluster a single instance of etcd can run on the same node with all the other master components.

However, for more substantial clusters it is typical to have a 3-node or even 5-node etcd cluster for redundancy and high availability.



kubernetes

# Controller manager



The controller-manager is a collection of various managers rolled up into one binary. It is the primary daemon that manages all core component control loops.

All these managers monitor the cluster state via the API and steers the cluster towards the desired state.

# Scheduler



The kube-scheduler is responsible for evaluating the workload requirements and attempts to place it on a matching resource.

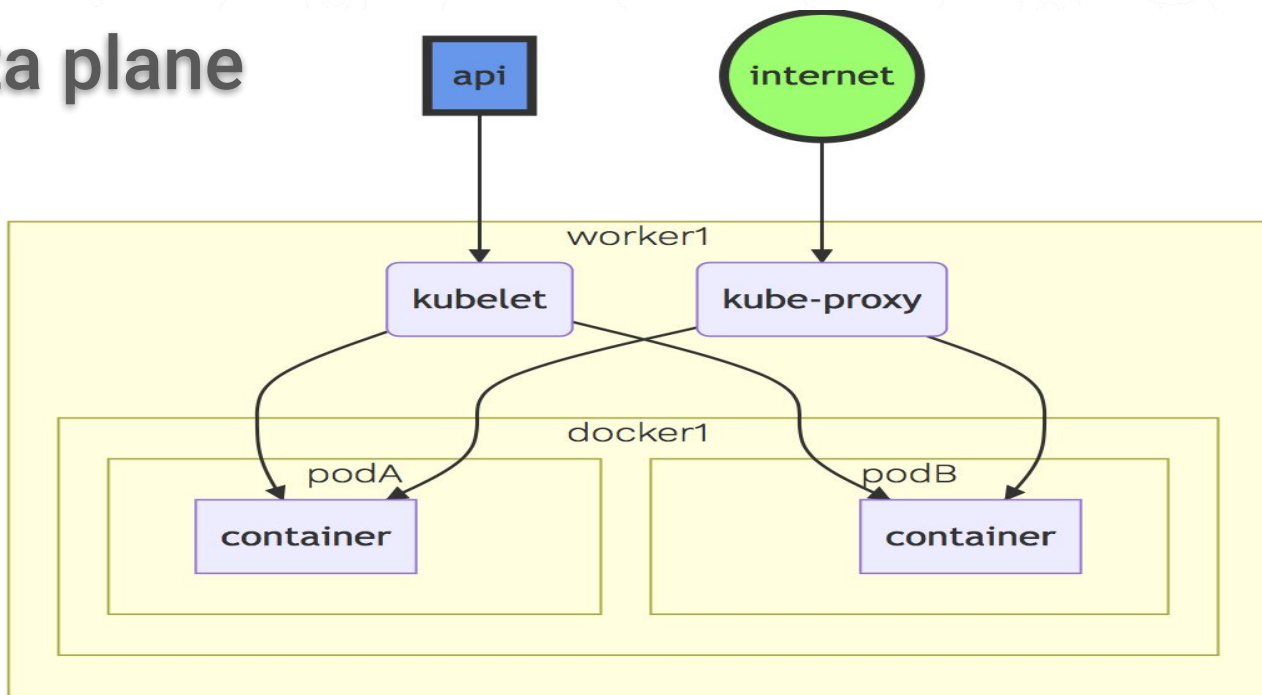
These requirements can include such things as general hardware reqs, affinity, anti-affinity, and other custom resource requirements.



# Node components (Data plane)



# Data plane



- Made up of worker nodes
- kubelet: Acts as a conduit between the API server and the node
- kube-proxy: Manages IP translation and routing



# kubelet



Acts as the node agent responsible for managing pod lifecycle on its host.

It oversees the communication with the master components and manage the running pods.

- Download pod secrets from the API server
- Mount volumes
- Run the pod's container
- Report the status of the node and each pod
- Run container liveness probes

# kube-proxy



Manages the network rules on each node and performs connection forwarding or load balancing for Kubernetes cluster services.

# Container runtime



A container runtime is a CRI (Container Runtime Interface) compatible application that executes and manages containers.

Kubernetes originally supported Docker as a container runtime engine. However, that is no longer the case;

- Containerd (docker)
- CRI-O
- Rkt
- etc




# Cluster DNS



Provides Cluster Wide DNS for Kubernetes Services.

- kube-dns (default pre 1.11)
- CoreDNS (current default 1.13)



<https://www.katacoda.com/smclernon/scenarios/kubernetes-101>



# Kubernetes objects



# Namespaces



Namespaces are a logical cluster or environment, and are the primary method of partitioning a cluster or scoping access.

```
apiVersion: v1
kind: Namespace
metadata:
  name: prod
  labels:
    app: CorpWebApp
```

```
$ kubectl get ns --show-labels
```

NAME	STATUS	AGE	LABELS
default	Active	11h	<none>
kube-public	Active	11h	<none>
kube-system	Active	11h	<none>
prod	Active	6s	app=CorpWebApp





# Default namespaces



- **default**: The default namespace for any object without a namespace.
- **kube-system**: Acts as the home for objects and resources created by Kubernetes itself.
- **kube-public**: A special namespace; readable by all users that is reserved for cluster bootstrapping and configuration.

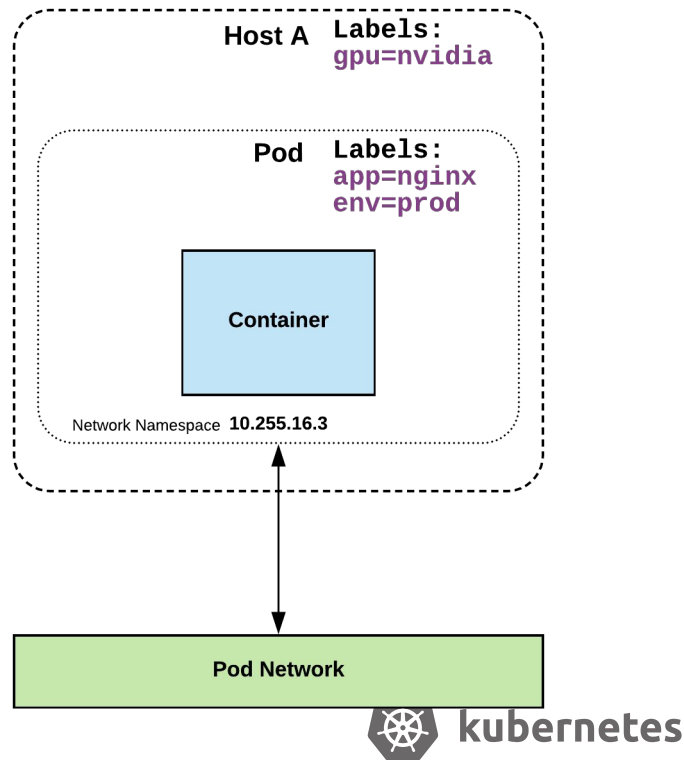
```
$ kubectl get ns --show-labels
```

NAME	STATUS	AGE	LABELS
default	Active	11h	<none>
kube-public	Active	11h	<none>
kube-system	Active	11h	<none>



# Labels


- key-value pairs that are used to identify, describe and group together related sets of objects or resources.
- **NOT** characteristic of uniqueness.
- Have a strict syntax with a slightly limited character set.



# Pods



- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- Foundational building block of Kubernetes Workloads.
- Pods are one or more containers that share volumes, a network namespace, and are a part of a **single context**.



# Pods are mortal



- They are born and when they die, they are not resurrected.
- ReplicaSets in particular create and destroy Pods dynamically (e.g. when scaling out or in).
- While each Pod gets its own IP address, even those IP addresses cannot be relied upon to be stable over time.

# Services



- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource** (unlike Pods)
  - static cluster-unique IP
  - static namespaced DNS name

**<service name>.<namespace>.svc.cluster.local**



## Services (contd)



- Target Pods using **equality based selectors**.
- Uses **kube-proxy** to provide simple load-balancing.
- **kube-proxy** acts as a daemon that creates **local entries** in the host's iptables for every service.



# Deployments



- Declarative method of managing Pods via **ReplicaSets**.
- Provide rollback functionality and update control.
- Updates are managed through the **pod-template-hash** label.
- Each iteration creates a unique label that is assigned to both the **ReplicaSet** and subsequent Pods.



# Deployment




- **revisionHistoryLimit**: The number of previous iterations of the Deployment to retain.
- **strategy**: Describes the method of updating the Pods based on the **type**. Valid options are **Recreate** or **RollingUpdate**.
  - **Recreate**: All existing Pods are killed before the new ones are created.
  - **RollingUpdate**: Cycles through updating the Pods according to the parameters: **maxSurge** and **maxUnavailable**.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    <pod template>
```







# And more



- **Ingress**
  - L7 load balancing
- **Jobs**
  - Run to completion
- **Autoscaling**
  - Automatically adjust replica count
- **DaemonSets**
  - Run something on every node (or subset)
- **StatefulSet**
  - Support for long term stateful distributed systems



# So much more



- **Role Based Access Control (RBAC)**
  - Control what users have access to what objects
- **Multiple Schedulers**
- **Flexible Scheduling Constraints**
  - Affinity, anti-affinity, taints, tolerations
- **Automatic Cluster Scaling**
  - K8s publishes signals that allow external services to scale the cluster automatically.
- **Cloud Provider Integration**
  - GCP, AWS, Azure, OpenStack, vSphere
- **Network Policy**
  - Network ingress policy



---

**question: 'Are we all  
YAML engineers now'**



# Gotchas



If you are managing your own cluster, then it is not without its perils.

Even with all the features that Kubernetes provides, it still requires lots of monitoring and pre-emptive capacity planning.

# Unbalanced clusters





# Unmonitored clusters will go on fire (not literally)



# Without autoscaling in place, clusters can crash







# Kubernetes Certified Service Provider


At time of writing there are 60+ certified service providers - see full list at <https://www.cncf.io/certification/kcsp/>



**Accenture (KCSP)** MCap: \$101B  
Accenture



**Alauda (KCSP)** Funding: \$15M  
Alauda



**Alibaba Cloud (KCSP)** MCap: \$390B  
Alibaba Cloud



**Banzai Cloud (KCSP)** Funding: \$2.5M  
Banzai Cloud



**Bitnami (KCSP)**  
Bitnami



**BoCloud (KCSP)** Funding: \$15.3M  
Bocloud



**BoxBoat (KCSP)**  
BoxBoat Technologies



**Caicloud (KCSP)** Funding: \$7.3M  
Caicloud



**Canonical (KCSP)**  
Canonical



**CloudOps (KCSP)**  
CloudOps



**CloudYuga (KCSP)**  
CloudYuga



**CloudZone (KCSP)** MCap: \$503M  
Matrix



**ComponentSoft**  
open cloud trainings and consultation

**Component Soft (KCSP)**  
Component Soft



**Container Solutions (KCSP)**  
Container Solutions



**CoreOS (KCSP)** MCap: \$31.2B  
Red Hat



**CREATIONLINE, INC.**

**Creationline (KCSP)**  
Creationline



**DaoCloud**

**DaoCloud (KCSP)** Funding: \$14.6M  
DaoCloud



**DATA ESSENTIAL**  
MAKE SENSE OF YOUR DATA

**Data Essential (KCSP)**  
Data Essential



**DoiT International (KCSP)**  
DoiT International



**EasyStack (KCSP)** Funding: \$110M  
EasyStack



**eKing Technology**  
易|建|科|技

**eKing Technology (KCSP)**  
Hainan eKing Technology



**Elastisys (KCSP)** Funding: \$680K  
Elastisys



**ELASTX (KCSP)**  
ELASTX



**Ghostcloud**

**Ghostcloud (KCSP)**  
Ghostcloud



**Giant Swarm (KCSP)** Funding: \$3.33M  
Giant Swarm



**Heptio (KCSP)** Funding: \$33.5M  
Heptio



**Huawei (KCSP)**  
Huawei Technologies



**IBM Cloud (KCSP)** MCap: \$110B  
IBM



**InfraCloud Technologies (KCSP)**  
InfraCloud Technologies



**inwinSTACK (KCSP)**  
inwinSTACK



Jetstack (KCSP)  
Jetstack



Kinvolk (KCSP)  
Kinvolk



Kublr (KCSP)  
Kublr



Kumina (KCSP)  
Kumina



LiveWyer (KCSP)  
LiveWyer



Loodse (KCSP)  
Loodse



MESOSPHERE

Mesosphere (KCSP) Funding: \$247M  
Mesosphere



MIRANTIS

Mirantis (KCSP) Funding: \$220M  
Mirantis



Mobilise Cloud Services (KCSP)  
Mobilise Cloud



MSys Technologies (KCSP)  
MSys Technologies



Nebulaworks (KCSP)  
Nebulaworks



nirmata

Nirmata (KCSP)  
Nirmata



OCTO Technology (KCSP)  
OCTO Technology



Oteemo (KCSP)  
Oteemo



Praqma (KCSP)  
Praqma



PRODYNA (KCSP)  
PRODYNA



Rackspace (KCSP) Funding: \$17.8M  
Rackspace



Rancher (KCSP) Funding: \$30M  
Rancher Labs



ReactiveOps (KCSP)  
ReactiveOps



RX-M (KCSP)  
RX-M



Samsung SDS (KCSP)  
Samsung SDS



SAP (KCSP) MCap: \$122B  
SAP



StackPointCloud (KCSP) MCap: \$16.2B  
NetApp



SUPERGIANT  
Supergiant (KCSP)  
Supergiant



TenxCloud (KCSP)  
TenxCloud



teutoStack (KCSP)  
teutoStack



Transwarp (KCSP) Funding: \$67.6M  
Transwarp



Treasure Data (KCSP) Funding: \$54M  
Treasure Data



Weaveworks (KCSP) Funding: \$20M  
Weaveworks



Wise2C (KCSP)  
Wise2C Technology




# Additional resources

# Links

- **Kubernetes documentation**  
<https://kubernetes.io/docs/home/>
- **Interactive Kubernetes Tutorials**  
<https://www.katacoda.com/courses/kubernetes>
- **Learn Kubernetes the Hard Way by Kelsey Hightower**  
<https://github.com/kelseyhightower/kubernetes-the-hard-way>
- **Awesome Kubernetes**  
<https://github.com/ramitsurana/awesome-kubernetes>



# Some useful books...



# Kubernetes cluster setup

There are many tools available to help bootstrap and configure a self-managed/commercial Kubernetes cluster.

Local development/learning;

- [Minikube](#)
- [Docker for Mac](#)
- [Docker for Windows](#)
- [Kubeadm](#)
- [katacoda.com/learn](https://katacoda.com/learn)

Commercial offerings;

- GKE - <https://cloud.google.com/kubernetes-engine/>
- EKS - <https://aws.amazon.com/eks/>
- AKS - <https://azure.microsoft.com/en-gb/services/kubernetes-service/>
- IKS - <https://www.ibm.com/cloud/container-service>
- PKS - <https://pivotal.io/platform/pivotal-container-service>
- [Kops](#)





